

CS Assignment 1 Readme

Mohammad Afroz Alam

1 Code Explanation

1.1 mlp.py

Contains a class MLP to build and train multi-layer perceptrons. It is used in the following way:

1. `init : mlp = MLP(nhidden_layers, nnodes, activation_function)`
 `nhidden_layers` : (int) number of layers
 `nnodes` : (list) number of nodes per layer
 `activation_function` : (str) non-linear function applied after affine transformation
2. `forward : mlp.forward(input_vector)`
 `input_vector` : (numpy array) a 1-D array of size 784

 returns the probability for each of the 10 classes
3. `cross_entropy_loss : mlp.cross_entropy_loss(target)`
 `target` : (int) integer storing the correct class

 returns cross entropy loss of our predicted model and target
4. `backward : mlp.backward(target)`
 `target` : (int)[optional] integer storing the class (if no argument is given, then target v

 returns derivative of loss with respect to input
5. `updateParams : mlp.updateParams(hyperParams, optimizer)`
 This function is used for updating parameters.
 Expects the following input
 `hyperParams` : (list) a list of hyperParams for the type of optimizer
 `optimizer` : (str) 'momentum' or 'adam'
 momentum:
 Takes the following params as a list:
 `learningRate` : hyperParams[0]
 `gamma` : hyperParams[1]

 adam:
 Takes the following params as a list:

```

        beta1 : hyperParams[0]
        beta2 : hyperParams[1]
        alpha : hyperParams[2]

6. plot_grads : mlp.plot_grads(fname, input_vector, target)
    fname      : (str) filename to save the plotted gradients
    input_vector : (array)[optional] input for forward pass
    target      : (int)[optional] for loss calculation

    this function plot a graph between numerical gradients and calculated gradients(using backpropagation)

7. train : mlp.train( train_data, train_labels, val_data, val_labels, config)
    train_data  : (array) a 2-D array with shape (trainsize, 784) to train MLP
    train_labels : (array) a 1-D array with shape (trainsize, ) to train MLP
    val_data     : (array) similar to train_data
    val_labels   : (array) similar to train_labels
    config : a dictionary with following entries
        filename : name of file to store results
        batchSize : batchSize for minibatch
        max_epochs : number of epochs to run
        optimizer : 'momentum' or 'adam'
        hyperParams:
            momentum:
                learningRate : hyperParams[0]
                gamma         : hyperParams[1]
            adam:
                beta1 : hyperParams[0]
                beta2 : hyperParams[1]
                alpha : hyperParams[2]

8. test : mlp.test( test_data, test_labels)
    test_data  : (array) a 2-D array with shape (testsize, 784) to train MLP
    test_labels : (array) a 1-D array with shape (testsize, ) to train MLP

```

1.2 main.py

Contains configurations for different models and used to run the classifier, plot gradients, etc.

2 Experiments

A total of 8 experiments have been performed using the following combinations of:

1. Number of hidden layers (2)[1, 2]
2. Activation function (2)[tanh and relu]
3. Optimazation algorithm (2)[momentum and adam]

Accuracy achieved on the testset in each of the above mentioned combination is listed below:

Number of Hidden Layers	Activation Function	Optimizer	Accuracy
1	Tanh	Adam	97.96%
	ReLU	Adam	97.44%
	Tanh	Momentum	94.28%
	ReLU	Momentum	95.80%
2	Tanh	Adam	97.39%
	ReLU	Adam	97.43%
	Tanh	Momentum	96.23%
	ReLU	Momentum	96.11%

Table 1: Accuracy for different architectures run for 10 epochs

2.1 Gradient Checking

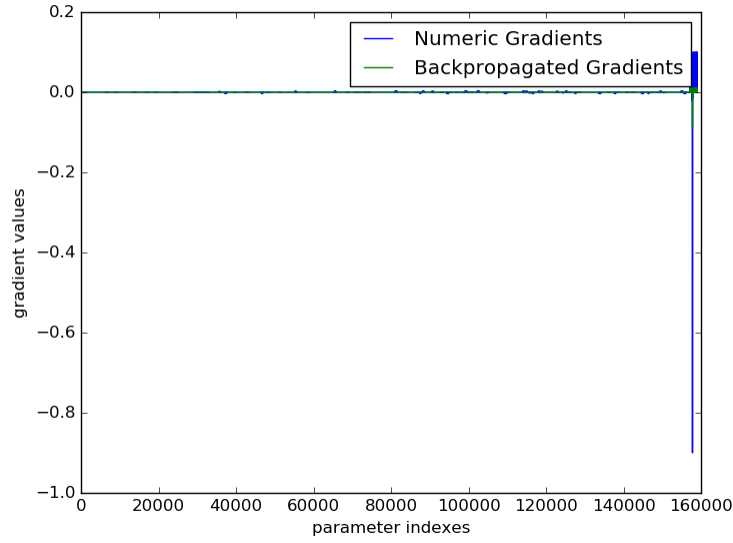


Figure 1: Backpropagated and Numerical Gradients

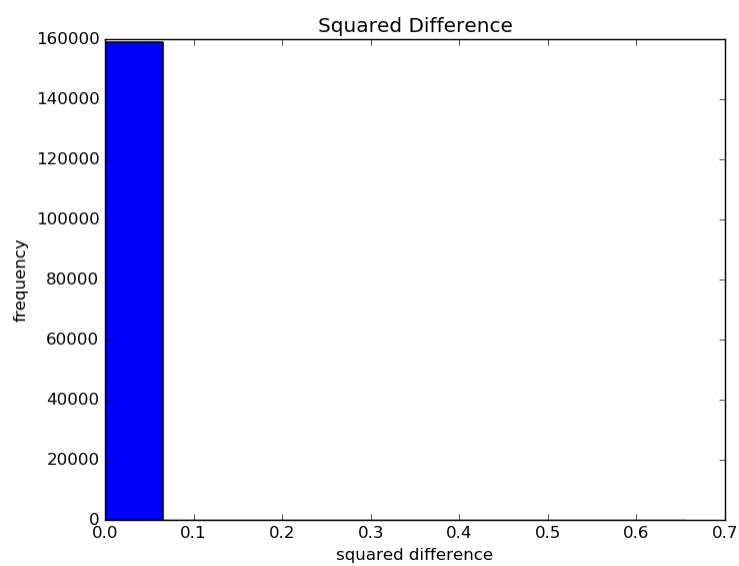


Figure 2: Histogram of squared error between Backpropagated and Numerical Gradients